

Iniciando o Elasticsearch em Docker

Docker Desktop Windows

```
PS E:\projetos\docker-elasticsearch\elastic> wsl -l -v
NAME                STATE              VERSION
* docker-desktop-data Running            2
  docker-desktop     Running            2
  Ubuntu-20.04        Running            2
```

```
PS E:\projetos\docker-elasticsearch\elastic> wsl -d docker-desktop
LAPTOP-V176DRSL:/tmp/docker-desktop-root/mnt/host/e/projetos/docker-elasticsearch/elastic# sysctl -w vm.max_map_count=262144
vm.max_map_count = 262144
LAPTOP-V176DRSL:/tmp/docker-desktop-root/mnt/host/e/projetos/docker-elasticsearch/elastic#
```

Docker Wsl2 Linux

```
feliciani@LAPTOP-V176DRSL:~$ sudo sysctl -w vm.max_map_count=262144
[sudo] password for feliciani:
vm.max_map_count = 262144
```

Docker Desktop Windows


```
PS E:\projetos\docker-elasticsearch\elastic> docker-compose up -d
Docker Compose is now in the Docker CLI, try 'docker compose up'




Starting elastic_elasticsearch_1 ... done
Starting elastic_kibana_1         ... done
Starting elastic_logstash_1       ... done
```

```
PS E:\projetos\docker-elasticsearch\elastic> docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
		NAMES			
d3d012693acc	docker.elastic.co/logstash/logstash:7.9.2	"/usr/local/bin/dock...	17 hours ago	Up 33 minutes	0.0.0.0:5044->5044/tcp,
:::5044->5044/tcp,	0.0.0.0:9600->9600/tcp, :::9600->9600/tcp	elastic_logstash_1			
ca700688aa0d	docker.elastic.co/kibana/kibana:7.9.2	"/usr/local/bin/dumb...	17 hours ago	Up 33 minutes	0.0.0.0:5601->5601/tcp,
:::5601->5601/tcp		elastic_kibana_1			
37a2fb5958f4	docker.elastic.co/elasticsearch/elasticsearch:7.9.2	"/tini -- /usr/local...	17 hours ago	Up 34 minutes	0.0.0.0:9200->9200/tcp,
:::9200->9200/tcp,	9300/tcp	elastic_elasticsearch_1			

Docker



 docker




Upgrade    Sign in

Containers / Apps


Images


Dev Environments


  elastic
E:\projetos\docker-elasticsearch\elastic

Open in Visual Studio Code   


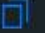
CONTAINERS

 elastic_logstash_1
docker.elastic.co/logstash/logstash:7.9.2
RUNNING PORT: 5044

 elastic_kibana_1
docker.elastic.co/kibana/kibana:7.9.2
RUNNING PORT: 5601

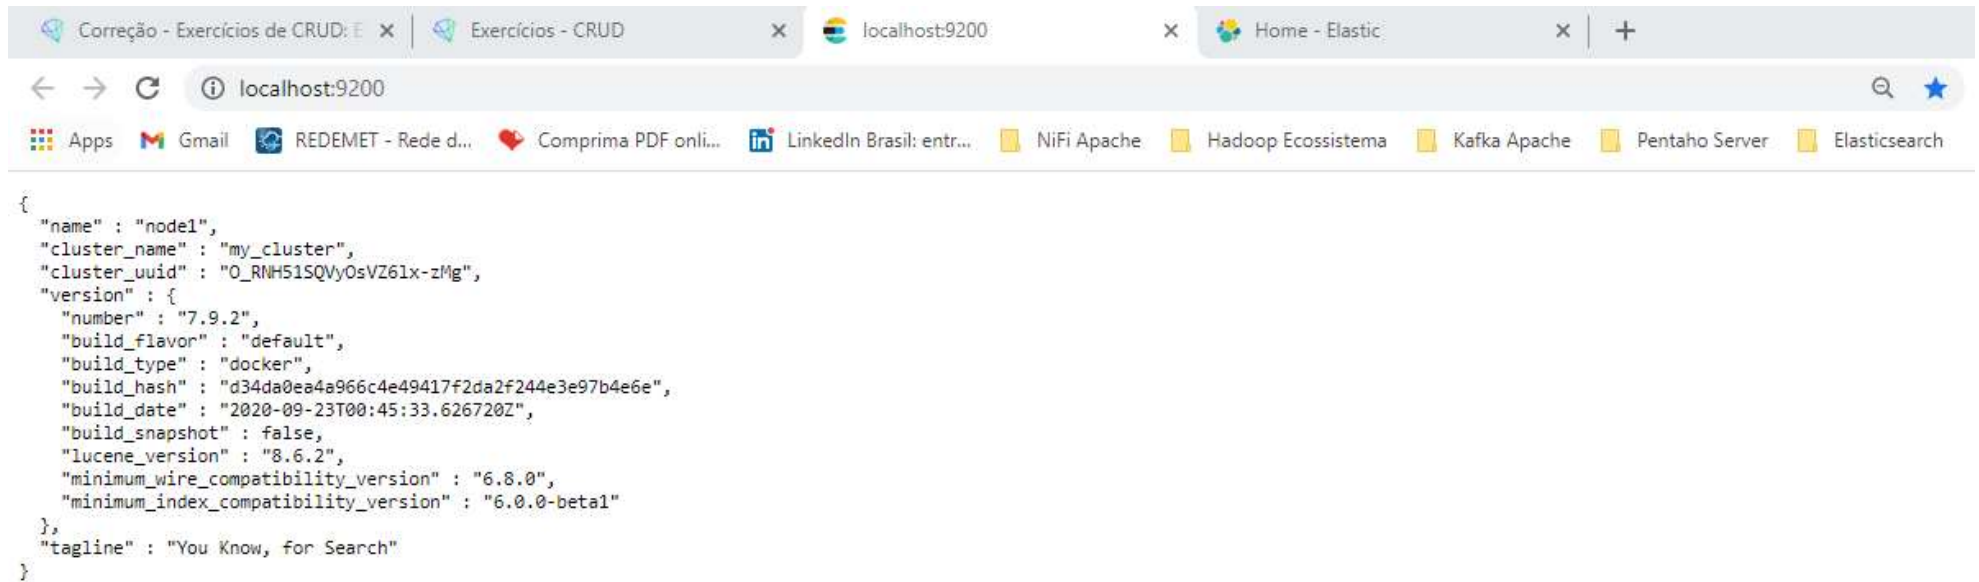
 elastic_elasticsearch_1
docker.elastic.co/elasticsearch/elasticsearch:7.9.2
RUNNING PORT: 9200

```
dest":"empty","referer":"http://localhost:5601/app/home","accept-encoding":"gzip, deflate, br","accept-language":"pt-BR,pt;q=0.9,en-US;q=0.8,en;q=0.7"},"remoteAddress":"172.18.0.1","userAgent":"Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/91.0.4472.77 Safari/537.36","referer":"http://localhost:5601/app/home"},"res":{"statusCode":200,"responseTime":751,"contentLength":9,"message":"POST /api/ui_metric/report 200 751ms - 9.0B"}
elasticsearch_1 | {"type": "server", "timestamp": "2021-06-10T13:42:53,810Z", "level": "INFO", "component": "o.e.m.j.JvmGcMonitorService", "cluster.name": "my_cluster", "node.name": "node1", "message": "[gc][1024] overhead, spent [304ms] collecting in the last [1s]", "cluster.uuid": "O_RNH51SQVyOsVZ6lx-zMg", "node.id": "I0MovYK2T0eNK1jQUwWviA" }
elasticsearch_1 | {"type": "server", "timestamp": "2021-06-10T13:50:02,818Z", "level": "WARN", "component": "o.e.m.f.FsHealthService", "cluster.name": "my_cluster", "node.name": "node1", "message": "health check of [/usr/share/elasticsearch/data/nodes/0] took [5569ms] which is above the warn threshold of [5s]", "cluster.uuid": "O_RNH51SQVyOsVZ6lx-zMg", "node.id": "I0MovYK2T0eNK1jQUwWviA" }
elasticsearch_1 | {"type": "server", "timestamp": "2021-06-10T13:50:12,644Z", "level": "WARN", "component": "o.e.m.j.JvmGcMonitorService", "cluster.name": "my_cluster", "node.name": "node1", "message": "[gc][young][1454][20] duration [2.5s], collections [1]/[3.1s], total [2.5s]/[4.5s], memory [367.7mb]->[86.2mb]/[512mb], all_pools {[young][282mb]->[0b]/[0b]}{[old][76.2mb]->[76.2mb]/[512mb]}{[survivor][9.4mb]->[10mb]/[0b]}", "cluster.uuid": "O_RNH51SQVyOsVZ6lx-zMg", "node.id": "I0MovYK2T0eNK1jQUwWviA" }
elasticsearch_1 | {"type": "server", "timestamp": "2021-06-10T13:50:12,645Z", "level": "WARN", "component": "o.e.m.j.JvmGcMonitorService", "cluster.name": "my_cluster", "node.name": "node1", "message": "[gc][1454] overhead, spent [2.5s] collecting in the last [3.1s]", "cluster.uuid": "O_RNH51SQVyOsVZ6lx-zMg", "node.id": "I0MovYK2T0eNK1jQUwWviA" }
```

Search...  Stick to bottom 

Acessado o Elasticsearch

<http://localhost:9200>



```
{
  "name" : "node1",
  "cluster_name" : "my_cluster",
  "cluster_uuid" : "O_RNH51SQVyOsVZ61x-zMg",
  "version" : {
    "number" : "7.9.2",
    "build_flavor" : "default",
    "build_type" : "docker",
    "build_hash" : "d34da0ea4a966c4e49417f2da2f244e3e97b4e6e",
    "build_date" : "2020-09-23T00:45:33.626720Z",
    "build_snapshot" : false,
    "lucene_version" : "8.6.2",
    "minimum_wire_compatibility_version" : "6.8.0",
    "minimum_index_compatibility_version" : "6.0.0-beta1"
  },
  "tagline" : "You Know, for Search"
}
```

Acessando o KIBANA

<http://localhost:5601>

The screenshot shows a web browser window with the Kibana interface. The browser's address bar displays `localhost:5601/app/home/`. The browser's tab bar shows several open tabs: "Correção - Exercícios de CRUD", "Exercícios - CRUD", "localhost:9200", and "Home - Elastic". The browser's bookmark bar contains links to "Apps", "Gmail", "REDEMET - Rede d...", "Comprima PDF onli...", "LinkedIn Brasil: entr...", "NiFi Apache", "Hadoop Ecosistema", "Kafka Apache", "Pentaho Server", and "Elasticsearch".

The Kibana interface features a left-hand navigation sidebar with the following sections:

- Home** (selected)
- Recently viewed**: No recently viewed items
- Kibana**: Discover, Dashboard, Canvas, Maps, Machine Learning, Visualize
- Enterprise Search**: App Search, Workplace Search
- Observability**: Overview, Logs

The main content area displays several cards and sections:

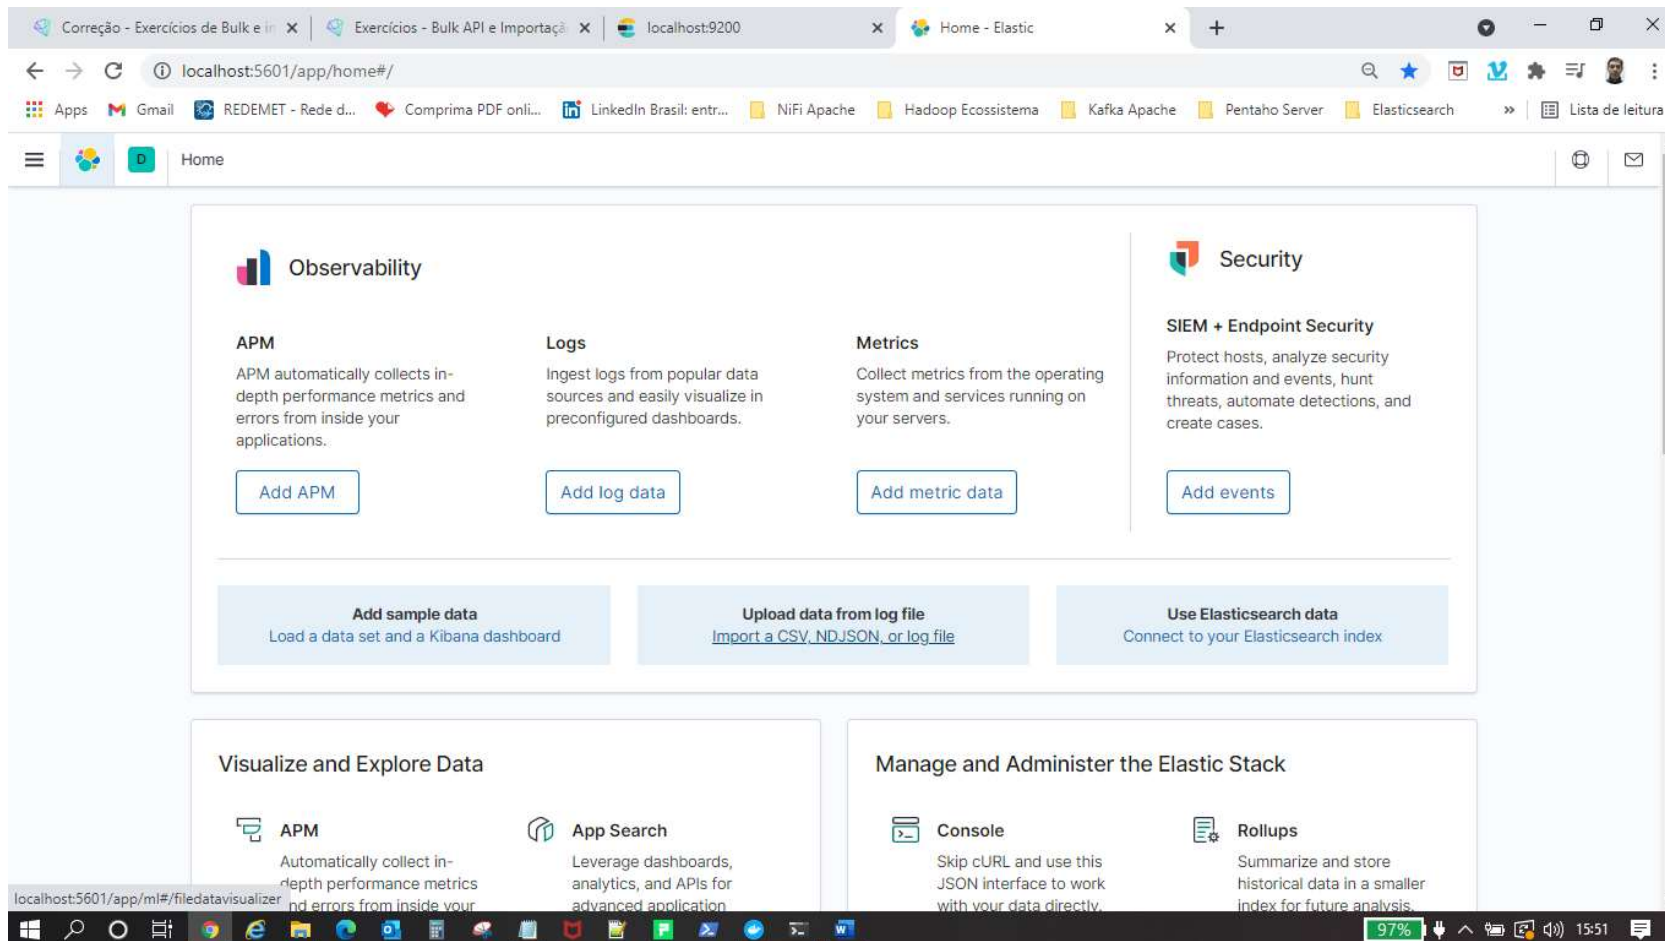
- Logs**: Ingest logs from popular data sources and easily visualize in preconfigured dashboards. Includes an "Add log data" button.
- Metrics**: Collect metrics from the operating system and services running on your servers. Includes an "Add metric data" button.
- Security**: SIEM + Endpoint Security. Protect hosts, analyze security information and events, hunt threats, automate detections, and create cases. Includes an "Add events" button.
- Data Ingestion**: Three buttons for "Add sample data", "Upload data from log file" (Import a CSV, NDJSON, or log file), and "Use Elasticsearch data" (Connect to your Elasticsearch index).
- Explore Data**: Includes a card for "App Search" (Leverage dashboards, analytics, and APIs for advanced application).
- Manage and Administer the Elastic Stack**: Includes cards for "Console" (Skip cURL and use this JSON interface to work with your data directly) and "Rollups" (Summarize and store historical data in a smaller interval for future analysis).

Exercitando a Bulk API e Importação

1. Importar os dados na Guia Arquivos para os índices

- Índice: concessionaria4
 - dataset/cars.bulk
- Índice: populacao
 - dataset/populacaoLA.csv

- Clicar em Upload data from log file



Clicar Select or drag and drop a file

Outra alternativa para chegar neste local: MENU -> Machine Learning -> Data Visualizer -> File

The screenshot shows a web browser window with the URL `localhost:5601/app/ml#/filedatavisualizer`. The browser's address bar and tabs are visible at the top. The page has a navigation bar with tabs: Overview, Anomaly Detection, Data Frame Analytics, Data Visualizer (selected), and Settings. The main content area is titled "Visualize data from a log file" with an "EXPERIMENTAL" badge. Below the title, there is a folder icon with a green plus sign. The text explains that the File Data Visualizer helps understand fields and metrics in a log file and allows uploading files up to 100 MB. It lists supported file formats: Delimited text files (CSV, TSV), Newline-delimited JSON, and Log files with a common timestamp format. A download icon is shown with the text "Select or drag and drop a file". A box at the bottom right indicates "Nenhum arquivo seleccionado". The Windows taskbar is visible at the bottom of the screen.

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Elastic x +

localhost:5601/app/ml#/filedatavisualizer

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecosystema Kafka Apache Pentaho Server Elasticsearch Lista de leitura

Machine Learning / Data Visualizer / File

Overview Anomaly Detection Data Frame Analytics **Data Visualizer** Settings

Visualize data from a log file EXPERIMENTAL

The File Data Visualizer helps you understand the fields and metrics in a log file. Upload your file, analyze its data, and then choose whether to import the data into an Elasticsearch index.

The File Data Visualizer supports these file formats:

- Delimited text files, such as CSV and TSV
- Newline-delimited JSON
- Log files with a common format for the timestamp

You can upload files up to 100 MB.

This feature is experimental. Got feedback? Please create an issue in [GitHub](#).

Select or drag and drop a file

Nenhum arquivo seleccionado

97% 15:55

Selecionei o arquivo cars.bulk para importar

The screenshot shows a web browser window with the URL `localhost:5601/app/ml#/filedatavisualizer`. The browser's address bar and tabs are visible at the top. Below the browser window, a file selection dialog titled "Abrir" is open, displaying a list of files in a directory named "05 - Elasticsearch Essential I > Dataset". The file `cars.bulk` is selected. The dialog also shows a sidebar with various file system locations and a bottom section with a "Nome" field and a file type dropdown.

File selection dialog details:

Nome	Data de modificação	Tipo	Tamanho
blogs.csv	07/06/2021 15:48	Arquivo de Valore...	5,200 KB
cars.bulk	07/06/2021 15:48	Arquivo BULK	1 KB
paris-925.logs	07/06/2021 15:49	Arquivo LOGS	66,059 KB
populacaoLA.csv	07/06/2021 15:48	Arquivo de Valore...	12 KB
weekly_MSFT.csv	07/06/2021 15:48	Arquivo de Valore...	59 KB

Below the dialog, a download icon and the text "Select or drag and drop a file" are visible.

Apresenta os dados que serão importados

The screenshot shows the Elastic ML interface in a web browser. The browser tabs include 'Correção - Exercícios de Bulk e...', 'Exercícios - Bulk API e Importação...', 'localhost:9200', and 'Elastic'. The address bar shows 'localhost:5601/app/ml#/filedatavisualizer'. The page title is 'Machine Learning / Data Visualizer / File'. The main content area is titled 'cars.bulk' and displays the 'File contents' section, showing the first 16 lines of a JSON file. Below this, a 'Summary' section provides statistics on the analyzed data.

File contents
First 16 lines

```
1 { "index": {} }
2 { "price" : 18000, "color" : "red", "make" : "honda", "sold" : "2016-10-28" }
3 { "index": {} }
4 { "price" : 20000, "color" : "red", "make" : "honda", "sold" : "2016-11-05" }
5 { "index": {} }
6 { "price" : 30000, "color" : "green", "make" : "ford", "sold" : "2016-05-18" }
7 { "index": {} }
8 { "price" : 15000, "color" : "blue", "make" : "toyota", "sold" : "2016-07-02" }
9 { "index": {} }
10 { "price" : 12000, "color" : "green", "make" : "toyota", "sold" : "2016-08-19" }
11 { "index": {} }
12 { "price" : 20000, "color" : "red", "make" : "honda", "sold" : "2016-11-05" }
13 { "index": {} }
14 { "price" : 80000, "color" : "red", "make" : "bmw", "sold" : "2016-01-02" }
```

Summary

Number of lines analyzed	16
Format	ndjson

At the bottom of the interface, there is a dark bar with two buttons: 'Import' and 'Cancel'. The Windows taskbar at the very bottom shows the system tray with a battery level of 97% and the time 16:02.

Clicar em Import

Correção - Exercícios de Bulk e in

Exercícios - Bulk API e Importaçã

localhost:9200

Elastic

localhost:5601/app/ml#/filedatavisualizer

AppsGmailREDEMET - Rede d...Comprima PDF onli...LinkedIn Brasil: entr...NiFi ApacheHadoop EcossistemaKafka ApachePentaho ServerElasticsearchLista de leitura

Machine Learning / Data Visualizer / File

14 { "price" : 80000, "color" : "red", "make" : "bmw", "sold" : "2016-01-02" }

15 { "index": {} }

16 { "price" : 25000, "color" : "blue", "make" : "ford", "sold" : "2016-02-12" }

Summary

Number of lines analyzed16

Formatndjson

Override settingsAnalysis explanation

File stats

color

8 documents (50%)

3 distinct values

top values

red50%

blue25%

green25%

make

8 documents (50%)

4 distinct values

top values

honda37.5%

ford25%

toyota25%

bmw12.5%

price

8 documents (50%)

7 distinct values

min10000

median20000

max80000

top values

2000025%

1000012.5%

1200012.5%

Import

Cancel

97%

16:02

Dar o nome de concessionaria4 e deixar marcada a opção Create index pattern para fazer Dashboard no Kibana com esses dados

The screenshot shows a web browser window with multiple tabs. The active tab is 'localhost:5601/app/ml#/filedatavisualizer'. The browser's address bar shows the URL. Below the browser window, the Kibana interface is visible. The top navigation bar includes 'Overview', 'Anomaly Detection', 'Data Frame Analytics', 'Data Visualizer' (which is selected), and 'Settings'. The main content area is titled 'cars.bulk'. Below this, there is a section for 'Import data' with a toggle for 'EXPERIMENTAL'. Underneath, there are two tabs: 'Simple' and 'Advanced'. The 'Simple' tab is active. In the 'Simple' tab, there is a field for 'Index name' containing the text 'concessionaria4'. Below this field, there is a checkbox labeled 'Create index pattern' which is checked. At the bottom of the 'Import data' section, there is a blue button labeled 'Import'. At the very bottom of the screen, there is a Windows taskbar with various application icons and a system tray showing the time as 16:16 and battery level at 97%.

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Elastic x +

localhost:5601/app/ml#/filedatavisualizer

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecossistema Kafka Apache Pentaho Server Elasticsearch » Lista de leitura

Machine Learning / Data Visualizer / File

Overview Anomaly Detection Data Frame Analytics **Data Visualizer** Settings

cars.bulk

Import data EXPERIMENTAL

Simple Advanced

Index name

concessionaria4

☒ Create index pattern

Import

Back Cancel

97% 16:16

Concessionaria4 foi importado

Correção - Exercícios de Bulk e inExercícios - Bulk API e Importaçãlocalhost:9200Elastic

localhost:5601/app/ml#/filedatavisualizer

AppsGmailREDEMET - Rede d...Comprima PDF onli...LinkedIn Brasil: entr...NiFi ApacheHadoop EcossistemaKafka ApachePentaho ServerElasticsearchLista de leitura

Machine Learning / Data Visualizer / File

Index name

concessionaria4

☒ Create index pattern

Reset

✓

File processed

✓

Index created

✓

Data uploaded

✓

Index pattern created

✓ Import complete

Indexconcessionaria4

Index patternconcessionaria4

Documents ingested16

BackCancel

97%

16:17

Inserindo outro Arquivo.

Import Data

The screenshot shows a web browser window with the URL `localhost:5601/app/ml#/datavisualizer`. The browser's address bar and tabs are visible at the top. The page has a navigation bar with tabs: Overview, Anomaly Detection, Data Frame Analytics, **Data Visualizer** (selected), and Settings. The main content area is titled "Data Visualizer" and includes a descriptive paragraph: "The Machine Learning Data Visualizer tool helps you understand your data, by analyzing the metrics and fields in a log file or an existing Elasticsearch index." Below this, there are three main interactive cards. The first card, labeled "EXPERIMENTAL", is titled "Import data" and contains the text "Import data from a log file. You can upload files up to 100 MB." with an "Upload file" button. The second card is titled "Select an index pattern" and contains the text "Visualize the data in an existing Elasticsearch index." with a "Select index" button. The third card, at the bottom, is titled "Start trial" and contains the text "To experience the full Machine Learning features that a [Platinum or Enterprise subscription](#) offers, start a 30-day trial." with a "Start trial" button. The browser's taskbar at the bottom shows various application icons and system status indicators, including a battery level of 97% and the time 16:19.

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Elastic x +

localhost:5601/app/ml#/datavisualizer

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecosystema Kafka Apache Pentaho Server Elasticsearch » Lista de leitura


Machine Learning / **Data Visualizer**

Overview Anomaly Detection Data Frame Analytics **Data Visualizer** Settings

Data Visualizer

The Machine Learning Data Visualizer tool helps you understand your data, by analyzing the metrics and fields in a log file or an existing Elasticsearch index.


EXPERIMENTAL



Import data

Import data from a log file. You can upload files up to 100 MB.

[Upload file](#)



Select an index pattern

Visualize the data in an existing Elasticsearch index.

[Select index](#)

Start trial

To experience the full Machine Learning features that a [Platinum or Enterprise subscription](#) offers, start a 30-day trial.

[Start trial](#)

localhost:5601/app/ml#/filedatavisualizer

97% 16:19

Select or drag and drop a file

The screenshot shows a web browser window with the URL `localhost:5601/app/ml#/filedatavisualizer`. The browser's address bar and tabs are visible at the top. The page has a navigation bar with tabs: Overview, Anomaly Detection, Data Frame Analytics, **Data Visualizer** (selected), and Settings. The main content area is titled "Visualize data from a log file" with an "EXPERIMENTAL" badge. Below the title, a text block explains the tool's purpose: "The File Data Visualizer helps you understand the fields and metrics in a log file. Upload your file, analyze its data, and then choose whether to import the data into an Elasticsearch index." To the left of this text is a folder icon with a green plus sign. Below the text, a list of supported file formats is provided: Delimited text files (CSV, TSV), Newline-delimited JSON, and Log files with a common timestamp format. Further down, it states "You can upload files up to 100 MB." and "This feature is experimental. Got feedback? Please create an issue in [GitHub](#)." At the bottom of the main content area, there is a large box with a download icon and the text "Select or drag and drop a file". Below this text is a button that says "Nenhum arquivo selecionado". The Windows taskbar is visible at the very bottom of the screen.

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Elastic x +

localhost:5601/app/ml#/filedatavisualizer

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecossistema Kafka Apache Pentaho Server Elasticsearch » Lista de leitura

Machine Learning / Data Visualizer / File

Overview Anomaly Detection Data Frame Analytics **Data Visualizer** Settings

Visualize data from a log file EXPERIMENTAL

The File Data Visualizer helps you understand the fields and metrics in a log file. Upload your file, analyze its data, and then choose whether to import the data into an Elasticsearch index.

The File Data Visualizer supports these file formats:

- Delimited text files, such as CSV and TSV
- Newline-delimited JSON
- Log files with a common format for the timestamp

You can upload files up to 100 MB.

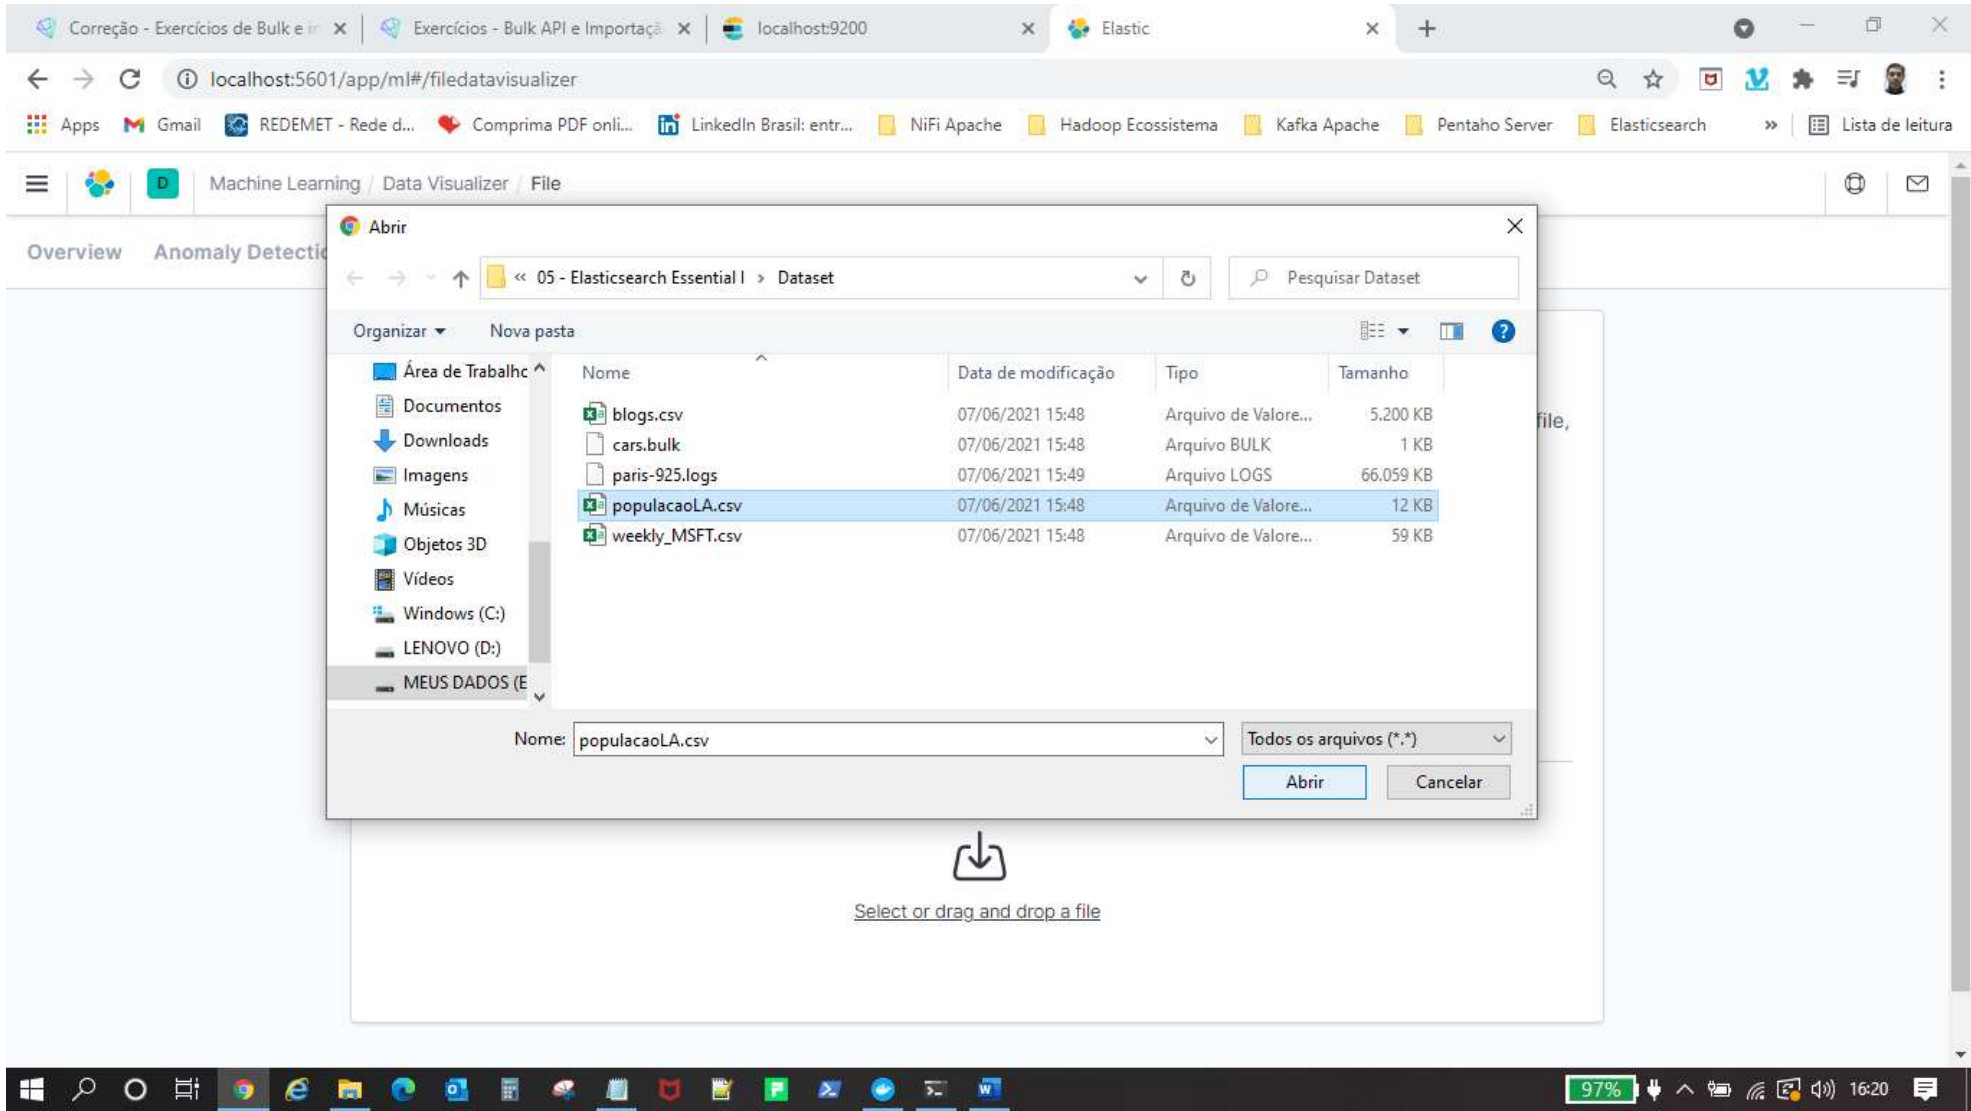
This feature is experimental. Got feedback? Please create an issue in [GitHub](#).

Select or drag and drop a file

Nenhum arquivo selecionado

97% 16:19

Arquivo populacaoLA.csv com dados da população de Los Angeles



Dados que serão importados

Correção - Exercícios de Bulk e inExercícios - Bulk API e Importaçãlocalhost:9200Elastic

localhost:5601/app/ml#/filedatavisualizer

AppsGmailREDEMET - Rede d...Comprima PDF onli...LinkedIn Brasil: entr...NiFi ApacheHadoop EcossistemaKafka ApachePentaho ServerElasticsearchLista de leitura

Machine Learning / Data Visualizer / File

OverviewAnomaly DetectionData Frame AnalyticsData VisualizerSettings

populacaoLA.csv

File contents

First 320 lines

1 Zip Code,Total Population,Median Age,Total Males,Total Females,Total Households,Average Household Size

2 91371,1,73.5,0,1,1,1

3 90001,57110,26.6,28468,28642,12971,4.4

4 90002,51223,25.5,24876,26347,11731,4.36

5 90003,66266,26.3,32631,33635,15642,4.22

6 90004,62100,34.8,31302,30878,22547,2.73

7 90005,37681,33.9,19299,18382,15044,2.5

8 90006,59185,32.4,30254,28931,18617,3.13

9 90007,40920,24,20915,20005,11944,3

10 90008,32327,39.7,14477,17850,13841,2.33

11 90010,3800,37.8,1874,1926,2014,1.87

12 90011,103892,26.2,52794,51098,22168,4.67

13 90012,31103,36.3,19493,11610,10327,2.12

14 90013,11772,44.6,7629,4143,6416,1.26

Summary

Number of lines analyzed

320

Format

delimited

Delimiter

,

ImportCancel

97%

16:22

Mostra algumas estatísticas, mesmo antes da importação

Import

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Elastic x +

localhost:5601/app/ml#/filedatavisualizer

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecossistema Kafka Apache Pentaho Server Elasticsearch » Lista de leitura

Machine Learning / Data Visualizer / File

Zip Code

319 documents (100%)

319 distinct values

min	median	max
90001	90807	93591

top values

90001	0.31%
90002	0.31%
90003	0.31%
90004	0.31%
90005	0.31%
90006	0.31%
90007	0.31%
90008	0.31%
90010	0.31%
90011	0.31%

Total Population

319 documents (100%)

313 distinct values

min	median	max
0	31481	105549

top values

0	1.88%
42399	0.63%
1	0.31%
3	0.31%
15	0.31%
117	0.31%
156	0.31%
328	0.31%
388	0.31%
476	0.31%

Median Age

319 documents (100%)

177 distinct values

min	median	max
0	37.1	74

top values

33.9	2.19%
0	1.88%
37.9	1.57%
41.2	1.57%
41.7	1.57%
42.7	1.57%
31.9	1.25%
32.4	1.25%
34	1.25%
38.2	1.25%

Total Males

319 documents (100%)

310 distinct values

min	median	max
0	15283	52794

Total Females

319 documents (100%)

313 distinct values

min	median	max
0	16202	53185

Total Households

319 documents (100%)

307 distinct values

min	median	max
0	10968	31087

Import Cancel

97% 16:25

Importação concluída com o nome de população e com o Create index pattern marcado para futuros Dashboards

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Elastic x +

localhost:5601/app/ml#/filedatavisualizer

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecossistema Kafka Apache Pentaho Server Elasticsearch » Lista de leitura

Machine Learning / Data Visualizer / File

Index name

populacao

☒ Create index pattern

Reset

File processed Index created Ingest pipeline created Data uploaded Index pattern created

✓ Import complete

Index	populacao
Index pattern	populacao
Ingest pipeline	populacao-pipeline
Documents ingested	319

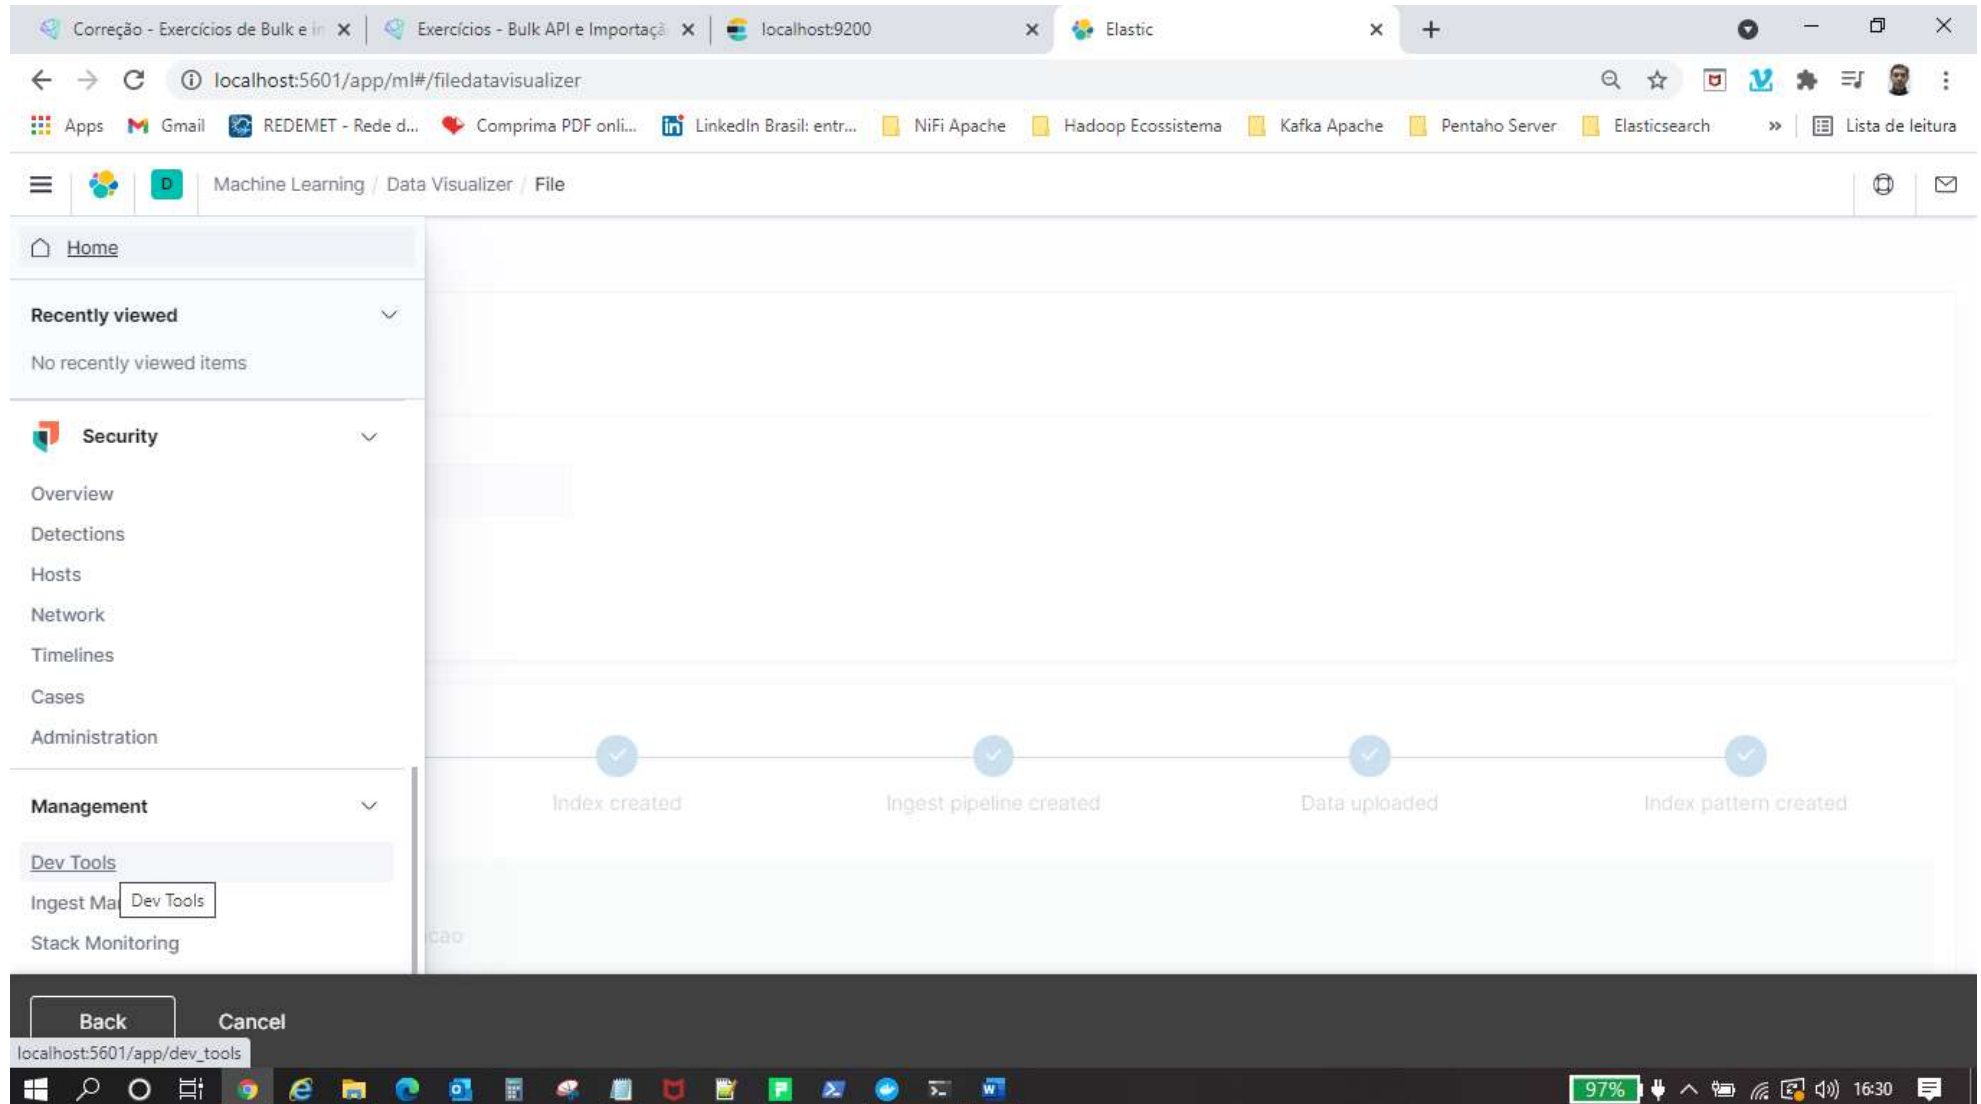
Back Cancel

97% 16:28

2. Executar as consultas

- Contar o número de documentos de cada um dos novos índices

Clicar em MENU e Dev Tools



Concessionaria4 contém 16 documentos

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Dev Tools - Elastic x +

localhost:5601/app/dev_tools#/console

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecossistema Kafka Apache Pentaho Server Elasticsearch » Lista de leitura

Dev Tools

Console Search Profiler Grok Debugger Painless Lab BETA

History Settings Help

200 - OK 13380 ms

```
1 GET concessionaria4/_count
2
3 GET produto/_count
4
5 GET produto/_doc/4
6
7 DELETE produto/_doc/4
8
9 GET produto/_doc/1
10
11 POST produto/_update/3
12 {
13   "doc": {
14     "qtd": 30
15   }
16 }
17
18 HEAD produto/_doc/3
19
20 POST produto/_doc/4
21 {
22   "nome": "CPU",
23   "qtd": 15,
24   "descricao": "15, 2.5Ghz"
25 }
26
27 POST produto/_doc/3
28 {
29   "nome": "memória ram",
30   "qtd": 10,
31   "descricao": "8GB, DDR4"
32 }
33
34 POST produto/_doc/4
```

```
1 {
2   "count" : 16,
3   "_shards" : {
4     "total" : 1,
5     "successful" : 1,
6     "skipped" : 0,
7     "failed" : 0
8   }
9 }
10
```

Windows taskbar: 97% battery, 16:32

População contém 319 documentos

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Dev Tools - Elastic x +

localhost:5601/app/dev_tools#/console

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecossistema Kafka Apache Pentaho Server Elasticsearch » Lista de leitura

Dev Tools

Console Search Profiler Grok Debugger Painless Lab BETA

History Settings Help

200 - OK 160 ms

```
1 GET populacao/_count
2
3 GET concessionaria4/_count
4
5 GET produto/_count
6
7 GET produto/_doc/4
8
9 DELETE produto/_doc/4
10
11 GET produto/_doc/1
12
13 POST produto/_update/3
14 {
15   "doc": {
16     "qtd": 30
17   }
18 }
19
20 HEAD produto/_doc/3
21
22 POST produto/_doc/4
23 {
24   "nome": "CPU",
25   "qtd": 15,
26   "descricao": "i5, 2.5Ghz"
27 }
28
29 POST produto/_doc/3
30 {
31   "nome": "memória ram",
32   "qtd": 10,
33   "descricao": "8GB, DDR4"
34 }
```

```
1 {
2   "count" : 319,
3   "_shards" : {
4     "total" : 1,
5     "successful" : 1,
6     "skipped" : 0,
7     "failed" : 0
8   }
9 }
10
```

97% 16:33

- Mostrar todos os documentos de cada um dos novos índices (16)

Mostra os 10 primeiros documentos da concessionaria4

The screenshot shows a web browser window with the DevTools console open. The console displays a series of REST client requests and responses. The first request is a GET to 'concessionaria4/_search', which returns a 200 OK status and a JSON response. The response indicates that 16 documents were found, with a total of 16 hits. The first two hits are shown in the expanded response view.

```
1 GET concessionaria4/_search
2
3 GET populacao/_count
4
5 GET concessionaria4/_count
6
7 GET produto/_count
8
9 GET produto/_doc/4
10
11 DELETE produto/_doc/4
12
13 GET produto/_doc/1
14
15 POST produto/_update/3
16 {
17   "doc": {
18     "qtd": 30
19   }
20 }
21
22 HEAD produto/_doc/3
23
24 POST produto/_doc/4
25 {
26   "nome": "CPU",
27   "qtd": 15,
28   "descricao": "i5, 2.5Ghz"
29 }
30
31 POST produto/_doc/3
32 {
33   "nome": "memória ram",
34   "qtd": 16
35 }
```

```
1 {
2   "took" : 38,
3   "timed_out" : false,
4   "_shards" : {
5     "total" : 1,
6     "successful" : 1,
7     "skipped" : 0,
8     "failed" : 0
9   },
10  "hits" : {
11    "total" : {
12      "value" : 16,
13      "relation" : "eq"
14    },
15    "max_score" : 1.0,
16    "hits" : [
17      {
18        "_index" : "concessionaria4",
19        "_type" : "doc",
20        "_id" : "IO5a93k8C_QioruG2f0j",
21        "_score" : 1.0,
22        "_source" : {
23          "index" : { }
24        }
25      },
26      {
27        "_index" : "concessionaria4",
28        "_type" : "doc",
29        "_id" : "Ie5a93k8C_QioruG2f0j",
30        "_score" : 1.0,
31        "_source" : {
32          "price" : 10000,
33          "color" : "red",
34          "..." : "..."
35        }
36      }
37    ]
38  }
39 }
```

200 - OK 170 ms

Mostra os 10 primeiros documentos da populacao (319)

Correção - Exercícios de Bulk e in x Exercícios - Bulk API e Importaçã x localhost:9200 x Dev Tools - Elastic x +

localhost:5601/app/dev_tools#/console

Apps Gmail REDEMET - Rede d... Comprima PDF onli... LinkedIn Brasil: entr... NiFi Apache Hadoop Ecosystema Kafka Apache Pentaho Server Elasticsearch » Lista de leitura

Dev Tools

Console Search Profiler Grok Debugger Painless Lab BETA

History Settings Help

200 - OK 120 ms

```
1 GET populacao/_search
2
3 GET concessionaria4/_search
4
5 GET populacao/_count
6
7 GET concessionaria4/_count
8
9 GET produto/_count
10
11 GET produto/_doc/4
12
13 DELETE produto/_doc/4
14
15 GET produto/_doc/1
16
17 POST produto/_update/3
18 {
19   "doc": {
20     "qtd": 30
21   }
22 }
23
24 HEAD produto/_doc/3
25
26 POST produto/_doc/4
27 {
28   "nome": "CPU",
29   "qtd": 15,
30   "descricao": "i5, 2.5Ghz"
31 }
32
33 POST produto/_doc/3
```

```
1 {
2   "took" : 10,
3   "timed_out" : false,
4   "_shards" : {
5     "total" : 1,
6     "successful" : 1,
7     "skipped" : 0,
8     "failed" : 0
9   },
10  "hits" : {
11    "total" : {
12      "value" : 319,
13      "relation" : "eq"
14    },
15    "max_score" : 1.0,
16    "hits" : [
17      {
18        "_index" : "populacao",
19        "_type" : "_doc",
20        "_id" : "Ie5k93k8C_QioruGMv4m",
21        "_score" : 1.0,
22        "_source" : {
23          "Total Population" : 1,
24          "Total Households" : 1,
25          "Total Females" : 1,
26          "Zip Code" : 91371,
27          "Median Age" : 73.5,
28          "Total Males" : 0,
29          "Average Household Size" : 1.0
30        }
31      },
32      {
33        "_index" : "populacao",
```